



DSBA Transformer survey paper study

A Survey of Transformers

#3: Attention 2

arXiv preprint



고려대학교 산업경영공학과

Data Science & Business Analytics Lab

이유경, 김명섭, 윤훈상, 김지나, 허재혁, 김수빈

발표자 : 김지나

1. Linearized Attention

1. Feature map
2. Aggregation rule

2. Prototype and Memory Compression

1. Attention with Prototype Queries
2. Attention with Compressed Key-Value Memory

1. Lineaerized Attention

- Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention (ICML 2020, 110회 인용)
- Masked Language Modeling for Proteins via Linearly Scalable Long-Context Transformers (arXiv 2020, 17회 인용)
- Random Feature Attention (ICLR 2021, 21회 인용)
- Rethinking Attention with Performers. (ICLR 2021, 117회 인용)
- Linear Transformers Are Secretly Fast Weight Memory Systems (arXiv 2021, 5회 인용)

2. Prototype and Memory Compression

- Fast Transformers with Clustered Attention (arXiv 2020, 20회 인용)
- Informer: Beyond Efficient Transformer for Long Sequence Time-Series Forecasting (AAAI 2021, 17회 인용)
- Generating Wikipedia by Summarizing Long Sequences(ICLR 2018, 379회 인용)
- Set Transformer (In Proceedings of ICML, 2019, 229회 인용), Luna (arXive 2021, 1회 인용)
- Luna: Linear Unified Nested Attention. (arXiv 2021, 1회 인용)

Linearization을 통해 attention의 computational complexity $\mathcal{O}(T^2) \rightarrow \mathcal{O}(T)$ 줄임

- 기존 Attention: $Q, K, V \in R^{T \times D}$ 에 대한 attention matrix를 위한 $\text{softmax}(QK^T)V$ 연산
 - QK^T 연산은 Quadratic, computational complexity $\mathcal{O}(T^2)$
- Linearized Attention: $\text{softmax}(QK^T)$ 연산을 위해 QK^T 를 $Q'K'^T$ 로 disentangle
 - Computational complexity $\mathcal{O}(T)$
 - K'^TV 연산 먼저 수행 후, Q' 와 연산
 - $Q'K'^TV \Rightarrow Q'(K'^TV)$

Linearized Attention

- Un-normalized attention matrix

$$\hat{\mathbf{A}} = \exp(\mathbf{Q}\mathbf{K}^\top)$$

- $\exp(\cdot)$ is applied element-wise
- 기존 softmax 취한 score에 따른 attention matrix에서 normalization을 위한 denominator 생략
- Regular Attention

$$\text{Attention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax}\left(\frac{\mathbf{Q}\mathbf{K}^\top}{\sqrt{D_k}}\right) \mathbf{V} = \mathbf{A}\mathbf{V}$$

$$\mathbf{Z} = \mathbf{D}^{-1} \hat{\mathbf{A}} \mathbf{V} \quad \text{where } \mathbf{D} = \text{diag}(\hat{\mathbf{A}} \mathbf{1}_T^\top)$$

$\mathbf{1}_T^\top$: the all-ones column vector of length T

Linearized Attention

- Approximate or replace the unnormalized attention matrix $\exp(QK^T)$ with $\phi(Q)\phi(K)^T$

$$\hat{A} = \exp(QK^T) \quad \longrightarrow \quad \phi(Q)\phi(K)^T$$

- ϕ : is a feature map that is applied in row-wise manner

$$z_i = \sum_j \frac{\text{sim}(\mathbf{q}_i, \mathbf{k}_j)}{\sum_{j'} \text{sim}(\mathbf{q}_i, \mathbf{k}_{j'})} \mathbf{v}_j, \quad \longrightarrow$$

Regular Attention의 $\text{sim}(\cdot, \cdot)$

: the exponential of inner product $\exp(\langle \cdot, \cdot \rangle)$

$$\begin{aligned} z_i &= \sum_j \frac{\phi(\mathbf{q}_i)\phi(\mathbf{k}_j)^T}{\sum_{j'} \phi(\mathbf{q}_i)\phi(\mathbf{k}_{j'})^T} \mathbf{v}_j \\ &= \frac{\phi(\mathbf{q}_i) \sum_j \phi(\mathbf{k}_j) \otimes \mathbf{v}_j}{\phi(\mathbf{q}_i) \sum_{j'} \phi(\mathbf{k}_{j'})^T}, \end{aligned}$$

$\text{sim}(\cdot, \cdot)$: a kernel function $K(x, y) = \phi(x)\phi(y)^T$

\otimes : outer product

Linearized Attention

$$z_i = \sum_j \frac{\text{sim}(\mathbf{q}_i, \mathbf{k}_j)}{\sum_{j'} \text{sim}(\mathbf{q}_i, \mathbf{k}_{j'})} \mathbf{v}_j,$$

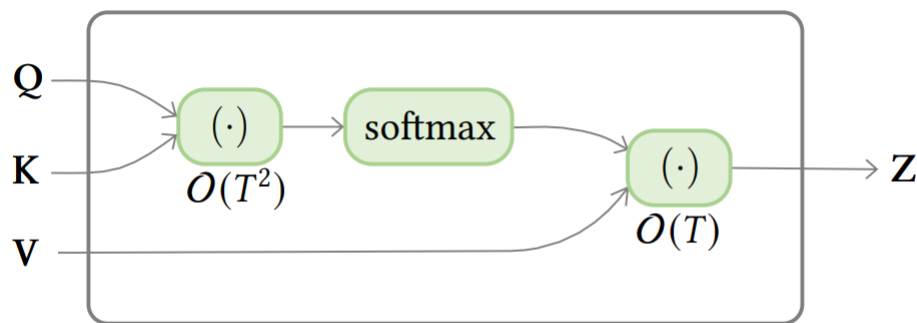
Regular Attention의 $\text{sim}(\cdot, \cdot)$

: the exponential of inner product $\exp(\langle \cdot, \cdot \rangle)$


$$\begin{aligned} z_i &= \sum_j \frac{\phi(\mathbf{q}_i) \phi(\mathbf{k}_j)^\top}{\sum_{j'} \phi(\mathbf{q}_i) \phi(\mathbf{k}_{j'})^\top} \mathbf{v}_j \\ &= \frac{\phi(\mathbf{q}_i) \sum_j \phi(\mathbf{k}_j) \otimes \mathbf{v}_j}{\phi(\mathbf{q}_i) \sum_{j'} \phi(\mathbf{k}_{j'})^\top}, \end{aligned}$$

$\text{sim}(\cdot, \cdot)$: a kernel function $K(x, y) = \phi(x) \phi(y)^\top$

\otimes : outer product



(a) standard self-attention




(b) linearized self-attention

Linearized Attention

$$\mathbf{z}_i = \sum_j \frac{\text{sim}(\mathbf{q}_i, \mathbf{k}_j)}{\sum_{j'} \text{sim}(\mathbf{q}_i, \mathbf{k}_{j'})} \mathbf{v}_j \quad \rightarrow \quad \mathbf{z}_i = \sum_j \frac{\phi(\mathbf{q}_i) \phi(\mathbf{k}_j)^\top}{\sum_{j'} \phi(\mathbf{q}_i) \phi(\mathbf{k}_{j'})^\top} \mathbf{v}_j$$

Regular Attention의 $\text{sim}(\cdot, \cdot)$

: the exponential of inner product $\exp(\langle \cdot, \cdot \rangle)$



Linearized Attention

$$z_i = \sum_j \frac{\text{sim}(\mathbf{q}_i, \mathbf{k}_j)}{\sum_{j'} \text{sim}(\mathbf{q}_i, \mathbf{k}_{j'})} \mathbf{v}_j,$$

Regular Attention의 $\text{sim}(\cdot, \cdot)$


: the exponential of inner product $\exp(\langle \cdot, \cdot \rangle)$

$\text{sim}(\cdot, \cdot)$: a kernel function $K(x, y) = \phi(x)\phi(y)^\top$

\otimes : outer product

$$\begin{aligned} z_i &= \sum_j \frac{\phi(\mathbf{q}_i)\phi(\mathbf{k}_j)^\top}{\sum_{j'} \phi(\mathbf{q}_i)\phi(\mathbf{k}_{j'})^\top} \mathbf{v}_j \\ &= \frac{\phi(\mathbf{q}_i) \sum_j \phi(\mathbf{k}_j) \otimes \mathbf{v}_j}{\phi(\mathbf{q}_i) \sum_{j'} \phi(\mathbf{k}_{j'})^\top}, \end{aligned}$$

Vector 연산의 summation



Linearized Attention

$$z_i = \sum_j \frac{\text{sim}(\mathbf{q}_i, \mathbf{k}_j)}{\sum_{j'} \text{sim}(\mathbf{q}_i, \mathbf{k}_{j'})} \mathbf{v}_j,$$

Regular Attention의 $\text{sim}(\cdot, \cdot)$

: the exponential of inner product $\exp(\langle \cdot, \cdot \rangle)$



$$\begin{aligned} z_i &= \sum_j \frac{\phi(\mathbf{q}_i) \phi(\mathbf{k}_j)^\top}{\sum_{j'} \phi(\mathbf{q}_i) \phi(\mathbf{k}_{j'})^\top} \mathbf{v}_j \\ &= \frac{\phi(\mathbf{q}_i) \sum_j \phi(\mathbf{k}_j) \otimes \mathbf{v}_j}{\phi(\mathbf{q}_i) \sum_{j'} \phi(\mathbf{k}_{j'})^\top}, \end{aligned}$$

Vector 연산의 summation

$\text{sim}(\cdot, \cdot)$: a kernel function $K(x, y) = \phi(x) \phi(y)^\top$

\otimes : outer product

Attention can be **linearized** by first computing the highlighted terms

연산량 매우 줄어듦

Linearized Attention

$$\begin{aligned}
 \mathbf{z}_i &= \sum_j \frac{\phi(\mathbf{q}_i)\phi(\mathbf{k}_j)^\top}{\sum_{j'} \phi(\mathbf{q}_i)\phi(\mathbf{k}_{j'})^\top} \mathbf{v}_j \\
 &= \frac{\phi(\mathbf{q}_i) \sum_j \phi(\mathbf{k}_j) \otimes \mathbf{v}_j}{\phi(\mathbf{q}_i) \sum_{j'} \phi(\mathbf{k}_{j'})^\top},
 \end{aligned}$$

- Memory matrix

$$\phi(\mathbf{q}_i) \sum_j \phi(\mathbf{k}_j) \otimes \mathbf{v}_j$$

- retrieve a value by multiplying the memory matrix with **feature mapped** query with proper normalization.

$$\sum_j \phi(\mathbf{k}_j) \otimes \mathbf{v}_j$$

- maintains a memory matrix **by aggregating associations** represented by outer products of (**feature mapped**) keys and values

(1) Feature map $\phi(\cdot)$ (2) Aggregation rule


Feature Maps

- Linear Transformer (ICML 2020, 110회 인용)
 - Simple feature map

$$\phi_i(\mathbf{x}) = \text{elu}(\mathbf{x}_i) + 1$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



기존의 dot product attention을 approximate하는 것을 목표로 하지 않고, 비슷한 수준의 성능을 내는 것을 목표로 하여, standard transformer에 준하는 성능 달성

Method	Validation PER	Time/epoch (s)
Bi-LSTM	10.94	1047
Softmax	5.12	2711
LSH-4	9.33	2250
Linear (ours)	8.08	824

- Speech recognition 실험 결과, linear transformer를 사용하였을 때, PER을 8까지 낮춰, 다른 모델에 비해 softmax와 가장 성능이 유사하며, 소요 시은 softmax의 3배 이상 감소했다.

Feature Maps

- Performer – first version (arXiv 2020, 17회 인용)
기존의 dot product attention을 approximate하는 것을 목표로 함
 - Random feature map

$$\phi(\mathbf{x}) = \frac{h(\mathbf{x})}{\sqrt{m}} [f_1(\omega_1^\top \mathbf{x}), \dots, f_m(\omega_m^\top \mathbf{x}), \dots, f_l(\omega_1^\top \mathbf{x}), \dots, f_l(\omega_m^\top \mathbf{x})],$$
$$f_1, \dots, f_l : \mathbb{R} \rightarrow \mathbb{R} \text{ and } h : \mathbb{R}^D \rightarrow \mathbb{R}.$$

where $\omega_1, \dots, \omega_m \stackrel{\text{iid}}{\sim} \mathcal{D}$ are drawn from some distribution $\mathcal{D} \in \mathcal{P}(\mathbb{R}^D)$

Softmax를 approximate하기 위해 아래와 같은 kernel 함수를 사용

$$h(\mathbf{x}) = \exp\left(\frac{\|\mathbf{x}\|^2}{2}\right), l = 2, f_1 = \sin, f_2 = \cos.$$

Feature Maps

- Performer – first version (arXiv 2020, 17회 인용)
기존의 dot product attention을 approximate하는 것을 목표로 함
 - Random feature map

Theorem 1 (Rahimi & Recht, 2007). *Let $\phi : \mathbb{R}^d \rightarrow \mathbb{R}^{2D}$ be a nonlinear transformation:*

$$\phi(\mathbf{x}) = \sqrt{1/D} \left[\sin(\mathbf{w}_1 \cdot \mathbf{x}), \dots, \sin(\mathbf{w}_D \cdot \mathbf{x}), \cos(\mathbf{w}_1 \cdot \mathbf{x}), \dots, \cos(\mathbf{w}_D \cdot \mathbf{x}) \right]^\top.$$

When d -dimensional random vectors \mathbf{w}_i are independently sampled from $\mathcal{N}(\mathbf{0}, \sigma^2 \mathbf{I}_d)$,

$$\mathbb{E}_{\mathbf{w}_i} [\phi(\mathbf{x}) \cdot \phi(\mathbf{y})] = \exp \left(-\|\mathbf{x} - \mathbf{y}\|^2 / 2\sigma^2 \right).$$

Feature Maps

- Random Feature Attention (ICLR 2021, 21회 인용)

Performer(ver. 1)와 유사

- Random feature map

$$\phi(\mathbf{x}) = \frac{h(\mathbf{x})}{\sqrt{m}} [f_1(\omega_1^\top \mathbf{x}), \dots, f_m(\omega_m^\top \mathbf{x}), \dots, f_l(\omega_1^\top \mathbf{x}), \dots, f_l(\omega_m^\top \mathbf{x})],$$
$$f_1, \dots, f_l : \mathbb{R} \rightarrow \mathbb{R} \text{ and } h : \mathbb{R}^D \rightarrow \mathbb{R}.$$

where $\omega_1, \dots, \omega_m \stackrel{\text{iid}}{\sim} \mathcal{D}$ are drawn from some distribution $\mathcal{D} \in \mathcal{P}(\mathbb{R}^D)$

query, key를 feature space에 보내기 전 \mathbf{l}_2 -normalization 하기 때문에, $h(\mathbf{x}) = \mathbf{1}$

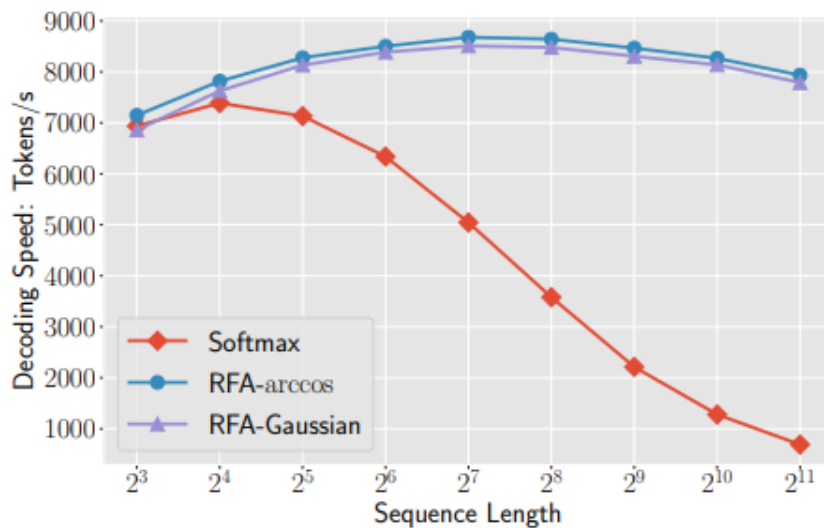
$$h(\mathbf{x}) = \exp\left(\frac{\|\mathbf{x}\|^2}{2}\right), l = 2, f_1 = \sin, f_2 = \cos.$$

Feature Maps


- Random Feature Attention (ICLR 2021, 21회 인용)

Performer(ver. 1)와 유사

- Random feature map



(a) Speed vs. lengths.



(b) Memory vs. lengths.

Feature Maps

- Random Feature Attention (ICLR 2021, 21회 인용)

Performer(ver. 1)와 유사

- Random feature map

The trigonometric random feature map leads to an unbiased approximation,

it **does not guarantee non-negative attention scores**

$$\phi(\mathbf{x}) = \frac{h(\mathbf{x})}{l} [f_1(\omega_1^\top \mathbf{x}), \dots, f_m(\omega_m^\top \mathbf{x}), \dots, f_l(\omega_1^\top \mathbf{x}), \dots, f_l(\omega_m^\top \mathbf{x})],$$

$$f_1, \dots, f_l : \mathbb{R} \rightarrow \mathbb{R} \text{ and } h : \mathbb{R}^D \rightarrow \mathbb{R}.$$

where $\omega_1, \dots, \omega_m \stackrel{\text{iid}}{\sim} \mathcal{D}$ are drawn from some distribution $\mathcal{D} \in \mathcal{P}(\mathbb{R}^D)$

query, key를 feature space에 보내기 전 \mathbf{l}_2 -normalization 하기 때문에, $h(\mathbf{x}) = 1$

$$h(\mathbf{x}) = \exp\left(\frac{\|\mathbf{x}\|^2}{2}\right), l = 2, f_1 = \sin, f_2 = \cos.$$

Feature Maps

- Performer – second version (ICLR 2021, 117회 인용)
 - Positive random feature map

$$\phi(\mathbf{x}) = \frac{h(\mathbf{x})}{\sqrt{m}} [f_1(\omega_1^\top \mathbf{x}), \dots, f_m(\omega_m^\top \mathbf{x}), \dots, f_l(\omega_1^\top \mathbf{x}), \dots, f_l(\omega_m^\top \mathbf{x})],$$
$$f_1, \dots, f_l : \mathbb{R} \rightarrow \mathbb{R} \text{ and } h : \mathbb{R}^D \rightarrow \mathbb{R}.$$

where $\omega_1, \dots, \omega_m \stackrel{\text{iid}}{\sim} \mathcal{D}$ are drawn from some distribution $\mathcal{D} \in \mathcal{P}(\mathbb{R}^D)$

$$h(\mathbf{x}) = \exp\left(-\frac{\|\mathbf{x}\|^2}{2}\right), l = 1, f_1 = \exp$$

Guarantees unbiased and non-negative approximation of dot-product attention
⇒ Performer(ver. 1) 보다 stable하고,
더 좋은 approximation 결과 보임

Feature Maps

- Performer – second version (ICLR 2021, 117회 인용)
 - Positive random feature map

$$\phi(\mathbf{x}) = \frac{h(\mathbf{x})}{\sqrt{m}} [f_1(\omega_1^\top \mathbf{x}), \dots, f_m(\omega_m^\top \mathbf{x}), \dots, f_l(\omega_1^\top \mathbf{x}), \dots, f_l(\omega_m^\top \mathbf{x})],$$
$$f_1, \dots, f_l : \mathbb{R} \rightarrow \mathbb{R} \text{ and } h : \mathbb{R}^D \rightarrow \mathbb{R}.$$

where $\omega_1, \dots, \omega_m \stackrel{\text{iid}}{\sim} \mathcal{D}$ are drawn from some distribution $\mathcal{D} \in \mathcal{P}(\mathbb{R}^D)$

$$h(\mathbf{x}) = 1, l = 1, f_1 = \text{ReLU}.$$

Effective in various tasks including machine translation and protein sequence modeling.

Feature Maps

- Linear Transformers Are Secretly Fast Weight Memory Systems (arXiv 2021, 5회 인용)
feature space 상에서의 orthogonality를 이용할 수 있는 feature map 제안

$$\phi_{i+2(j-1)D}(\mathbf{x}) = \text{ReLU}([\mathbf{x}, -\mathbf{x}])_i \text{ReLU}([\mathbf{x}, -\mathbf{x}])_{i+j}$$

for $i = 1, \dots, 2D, j = 1, \dots, v$.

input $x \in R^D$

the feature map $\phi : R^D \rightarrow R^{2vD}$

Aggregation Rule

The associations $\{\phi(k)_j \otimes v_j\}$ are aggregated into the memory matrix by simple summation

⇒ 새로운 association을 memory network S 에 추가할 때, 선택적으로 association을 drop하는 것이 효과적

Aggregation Rule

- Random Feature Attention (ICLR 2021, 21회 인용)
 - Gating mechanism

$$\begin{aligned}g_t &= \text{sigmoid}(\mathbf{w}_g \cdot \mathbf{x}_t + b_g), \\ \mathbf{S}_t &= g_t \mathbf{S}_{t-1} + (1 - g_t) \phi(\mathbf{k}_t) \otimes \mathbf{v}_t, \\ \mathbf{z}_t &= g_t \mathbf{z}_{t-1} + (1 - g_t) \phi(\mathbf{k}_t).\end{aligned}$$

- w_g and b_g are learned parameters, and \mathbf{x}_t is the input representation at timestep t .
- By multiplying the learned scalar gates $0 < g_t < 1$ against the hidden state $(\mathbf{S}_t, \mathbf{z}_t)$, history is exponentially decayed, favoring more recent context.

Aggregation Rule

- Linear Transformers Are Secretly Fast Weight Memory Systems (arXiv 2021, 5회 인용)

Association의 단순 합으로 memory matrix를 update하는 것은, memory matrix의 capacity를 제한하는 것, 따라서 write-and-remove를 통해 capacity를 확장하는 방식을 제안

- Write-and-remove update

$$\begin{aligned}
 \mathbf{k}^{(i)}, \mathbf{v}^{(i)}, \mathbf{q}^{(i)} &= \mathbf{W}_k \mathbf{x}^{(i)}, \mathbf{W}_v \mathbf{x}^{(i)}, \mathbf{W}_q \mathbf{x}^{(i)} \\
 \bar{\mathbf{v}}^{(i)} &= \mathbf{W}^{(i-1)} \phi(\mathbf{k}^{(i)}) \\
 \beta^{(i)} &= \sigma(\mathbf{W}_\beta \mathbf{x}^{(i)}) \\
 \mathbf{v}_{\text{new}}^{(i)} &= \beta^{(i)} \mathbf{v}^{(i)} + (1 - \beta^{(i)}) \bar{\mathbf{v}}^{(i)} \\
 \mathbf{W}^{(i)} &= \mathbf{W}^{(i-1)} + \underbrace{\mathbf{v}_{\text{new}}^{(i)} \otimes \phi(\mathbf{k}^{(i)})}_{\text{write}} - \underbrace{\bar{\mathbf{v}}^{(i)} \otimes \phi(\mathbf{k}^{(i)})}_{\text{remove}} \\
 &= \mathbf{W}^{(i-1)} + \beta^{(i)} (\mathbf{v}^{(i)} - \bar{\mathbf{v}}^{(i)}) \otimes \phi(\mathbf{k}^{(i)}) \\
 \mathbf{y}^{(i)} &= \mathbf{W}^{(i)} \phi(\mathbf{q}^{(i)})
 \end{aligned}$$

$\beta(i)$ is the “write-strength. only depends on input $x(i)$

새로운 input key-value pair가 들어오면,

- 1) k^i 와 직전 memory matrix W^{i-1} 의 association \bar{v}^i 을 구하고
- 2) 현재 v^i 와 \bar{v}^i 의 convex combination을 memory matrix에 update함

03 Query Prototyping and Memory Compression

개념

목적

Reduce the complexity of attention by

- (1) Query prototyping: reducing the number of queries
- (2) Memory compression: reducing the number of key-value pairs

The key-value pairs are often referred to as a **key-value memory**


03 Query Prototyping and Memory Compression

(1) Attention with Prototype Queries

Attention with Prototype Queries

Several prototypes of queries serve as the main source to compute attention distributions
query representation 중,

- 특정 position의 query distribution을 copy
- discrete uniform distribution으로 채움



(a) Query prototyping

03 Query Prototyping and Memory Compression

(1) Attention with Prototype Queries

Attention with Prototype Queries

- Clustered Attention (arXiv 2020, 20회 인용)
 - groups queries into several clusters and then computes attention distributions for cluster centroids.

① Centroid 구하기

$$Q_j^c = \frac{\sum_{i=1}^N S_{ij} Q_i}{\sum_{i=1}^N S_{ij}}$$

$S_{ij} = 1$, if the i -th query Q_i belongs to the j -th cluster and 0 otherwise.

② Centroid에 대한 attention score 계산 ③ Centroid에 의한 새로운 value 계산

$$A^c = \text{softmax} \left(\frac{Q^c K^T}{\sqrt{D_k}} \right)$$

$Q^c \in \mathbb{R}^{C \times D_k}$ as the centroid matrix

$$\hat{V}^c = A^c V.$$

④ 가장 가까운 centroid에 대한 attention value 도출

$$\hat{V}_i = \sum_{j=1}^C S_{ij} \hat{V}_j^c.$$

03 Query Prototyping and Memory Compression

(1) Attention with Prototype Queries

Attention with Prototype Queries

- Clustered Attention (arXiv 2020, 20회 인용)
 - groups queries into several clusters and then computes attention distributions for cluster centroids.

	full	clustered-100	i-clustered-100
WER (%)	15.0	18.5	15.5
Time/Epoch (h)	3.84	1.91	2.57
Convergence Time (h)	228.05	132.13	127.44

03 Query Prototyping and Memory Compression

(1) Attention with Prototype Queries

Attention with Prototype Queries


- Informer (AAAI 2021, 17회 인용)
 - Query sparsity measurement를 제안하여, 상위 u 개의 query만을 가지고 attention distribution 계산
 - 나머지 query에는 discrete uniform distribution 부여
 - Query sparsity measurement
 - Query의 attention distribution과 the discrete uniform distribution 사이의 Kullback-Leibler divergence 값을 기반으로 정의
 - Attention distribution이 the discrete uniform distribution과의 차이가 클수록 몇몇 key에 dominant한 attention을 주는 query라고 할 수 있기 때문에, KLD가 클수록 중요한 query로 봄

03 Query Prototyping and Memory Compression

(2) Attention with Compressed Key-Value Memory

Attention with Compressed Key-Value Memory

reduce the complexity by reducing the number of the key-value pairs before applying the attention mechanism



03 Query Prototyping and Memory Compression

(2) Attention with Compressed Key-Value Memory

Attention with Compressed Key-Value Memory


- Memory Compressed Attention(ICLR 2018, 379회 인용)
 - Strided convolution을 사용하여 key와 value의 개수를 줄임
 - Kernel size k 에 따라 줄이고자 하는 key, value 개수를 조절하며, attention 연산량이 줄어들기 때문에, 같은 시간 안에 vanilla transformer보다 훨씬 긴 sequence를 처리할 수 있다.
- Set Transformer (In Proceedings of ICML 2019, 229회 인용),
- Luna (arXive 2021, 1회 인용)
 - Trainable global node를 정의하여, input으로부터의 정보를 취합하게 하며, 취합된 정보는 input이 attend할 memory로 사용된다.

03 Query Prototyping and Memory Compression

(2) Attention with Compressed Key-Value Memory

Attention with Compressed Key-Value Memory

- Linformer(arXiv 2020, 132회 인용)
 - Key, value에 대해 linear projection을 적용하여 length n 에서 더 짧은 n_k 의 sequence로 사영
 - 단점: autoregressive attention에 사용할 수 없음
- Poolingformer (ICML 2021, 1회 인용)
 - 여러 개의 pooling operation을 사용하여, key와 value를 줄이면서도 receptive field를 키우고자 함



(a) Single-level local attention (b) Two-level pooling attention

1. Lineaerized Attention

- Transformers are RNNs: Fast Autoregressive Transformers with Linear Attention (ICML 2020, 110회 인용)
- Masked Language Modeling for Proteins via Linearly Scalable Long-Context Transformers (arXiv 2020, 17회 인용)
- Random Feature Attention (ICLR 2021, 21회 인용)
- Rethinking Attention with Performers. (ICLR 2021, 117회 인용)
- Linear Transformers Are Secretly Fast Weight Memory Systems (arXiv 2021, 5회 인용)

2. Prototype and Memory Compression

- Fast Transformers with Clustered Attention (arXiv 2020, 20회 인용)
- Informer: Beyond Efficient Transformer for Long Sequence Time-Series Forecasting (AAAI 2021, 17회 인용)
- Generating Wikipedia by Summarizing Long Sequences(ICLR 2018, 379회 인용)
- Set Transformer (In Proceedings of ICML, 2019, 229회 인용), Luna (arXiv 2021, 1회 인용)
- Luna: Linear Unified Nested Attention. (arXiv 2021, 1회 인용)

감사합니다